Brew What Sells with AI

An evidence engine for craft breweries

Ankur Napa   Growth Officer, Disruptive Advantage
Biergarten, Koramangala, Bengaluru   12 August 2026

Ankur Napa

Growth Officer, Disruptive Advantage

Eight years an R&D Master Brewer at AB InBev, then Principal Data Analyst for commercial planning there.

Beverage operations and digital transformation across beer, whisky and wine at iWort. Now AI and data for beverage companies.

I learned what a brewery decides before I learned how to model it. That order is why this session works.

HOW THIS HOUR RUNS

You answer first. Then I do.

01

A question goes up

One question on the screen, and no answer under it yet.

>

02

You scan the code

Phone camera on the code in the corner. No app and no login.

>

03

You type one line

Your own words, anonymous, and it is on the screen in seconds.

>

04

I read the room back

Your answers first, then mine. Where the two differ is the session.

Then the next question, and round again, spread through the hour.

You are not watching a demo. You are the input to it.

BEFORE I GIVE YOU MINE

What is AI?

Scan it. One line, in the words you would use on the floor.

No wrong answers. I will read some of them out.

Every answer in this room will be right about something and thin about the same thing. That gap is the next two slides.

WHICH PART OF AI

We are only in the smallest ring.

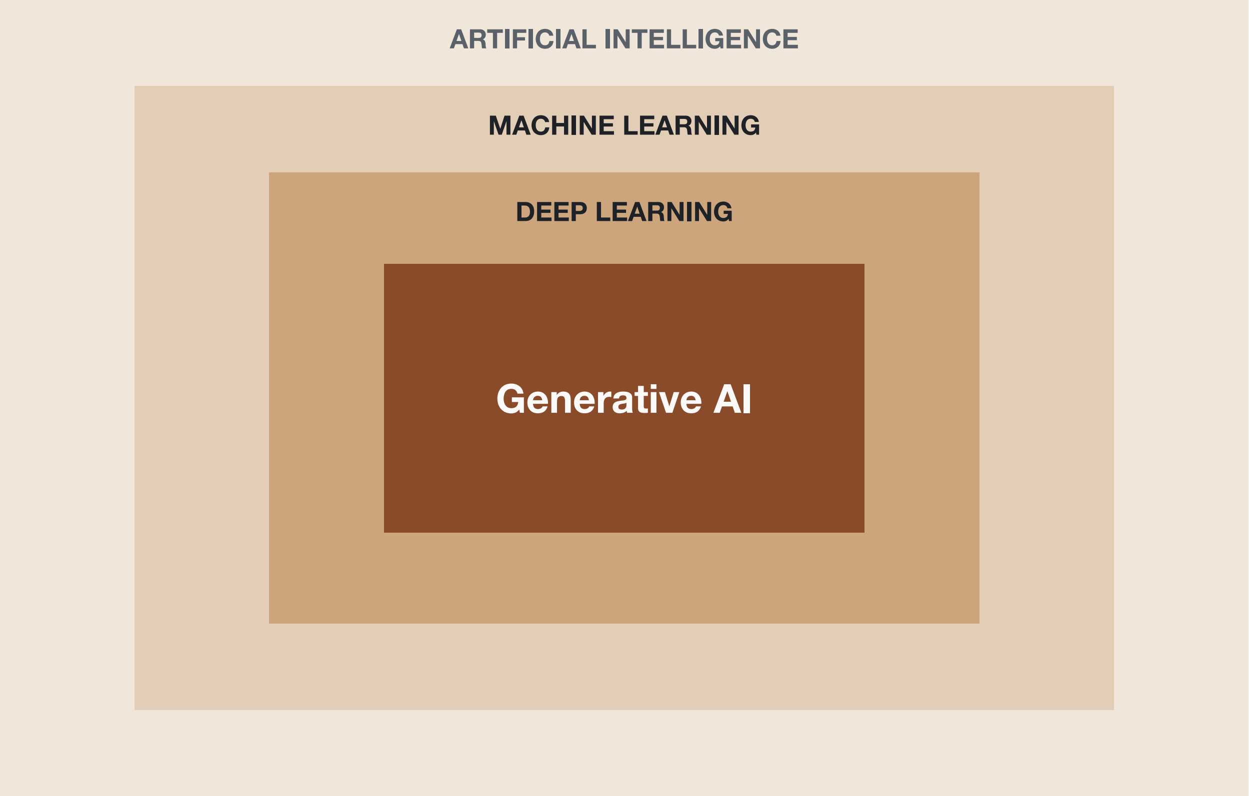

A chat model that reads the files you attach.

No sensors on your tanks. No model trained on your data. No project for anybody's IT department.

The outer rings are real, and they are what the trade press means by AI. None of them is what you will do on Wednesday.

APPS AND SERVICES IN THE MARKET

The elephant in the room.

ChatGPT	Gemini	**Claude**	Copilot
Perplexity	Llama	Grok	DeepSeek
Mistral	Midjourney	Notion AI	Gamma

And that is a partial list from this month.

Twelve logos, and the argument about which one is best is the loudest argument in the room and the least useful one.

They are the same kind of tool. The method is what transfers, and today it runs on one of them.

SECOND QUESTION

Paid or free?

How many of you already use one of these, in any form?

And is it the free version or a paid one?

Free will run most of what you see today. I will say clearly where it will not, and there are only two places.

TODAY'S HERO

Claude, and here is why.

1 **It reads what you attach, closely.**

Fifteen files and eighteen months of records in one thread, and it still has the first file when you get to the last question.

2 **It says a number is thin instead of inventing one.**

The honesty clause on your prompt card only works if the tool honours it. This is the one that most reliably does.

3 **It answers in plain language, not code.**

You ask for a table, you get a table. Nothing today needs a developer, and nothing today needs you to become one.

Pick a different one tomorrow and every prompt in your take-home pack still works.

WHICH CLAUDE

Three doors in. One of them is not today.

The Claude app

PHONE OR BROWSER

Attach a file, ask, read. Where most of you will start, probably in the car park afterwards.

Claude desktop

YOUR LAPTOP

The same thing, sitting closer to your files. Best for the monthly planning job with the exports open.

Claude Cowork

THE LONGER JOBS

Point it at a folder of exports and let it work through them while you do something else.

Claude Code

For developers, in a terminal. It is the most powerful of the four and it is not this hour. Nothing today needs it, and I will not open it.

THE BREWERY FOR TODAY

Deccan Draught Co.

Pune	Year 4	8	18 months
LOCATION	AGE	TAPS	OF RECORDS

AND THE FILES IT ALREADY EXPORTS

- Brewing process data
- Inventory
- Sales of beer
- POS data
- Social media

Fifteen files. Every one of them is something you can export this afternoon from what you already run.

The brewery is not real. Every number in it is, built from Brewers Association benchmarks and 250 real recipes.

Before you paste anything real.

Five minutes that will save somebody in this room a bad week.

DATA SECURITY

Do, and do not.

DO	DO NOT
Strip customer names, phone numbers and staff records first. Quantities and prices are fine.	Paste bank details, GST logins, or a contract with a confidentiality clause.
Turn off the setting that lets your chats train the model, on any account holding real numbers.	Use one shared brewery login. You will never know who uploaded what.
Keep one thread per decision, and close it when the decision is made.	Assume a deleted chat is gone everywhere the moment you delete it.
Read the output before it goes near a filing or a compliance document.	Put your recipe IP into a shared workspace you do not control.

THE ONE TEST

Before the file goes up.

Would you be comfortable if this file turned up in a competitor's inbox?

If not, take out the column that makes the answer no. It is almost always a name, and almost never a number.

Your sales quantities are commercially dull on their own. Your customer list is not. Send the first, keep the second.

THIRD QUESTION

What would you ask it?

You now have a brewery, eighteen months of its records, and the rules about what is safe to hand over.

So what is the first question you would put to it?

Whatever you type, somebody in this room has the same problem and has been solving it by hand for years.

FIVE SEATS

The same data. Five different questions.

01

Brewery owner

What makes money, and what only looks like it does?

02

Brewer

What do I brew next, and will the tanks take it?

03

Brewing incharge

Can we actually run this plan on the floor next month?

04

Social media handler

What do I post, and what am I not allowed to post?

05

Inventory manager

What do I order, and when does the lead time bite?

Everything today is inventory, sales and POS questions, asked from one of these five seats.

Pick yours now. The three use cases each belong to one of them.

WHAT YOU WILL WITNESS

Three things you leave with.

1. **Rank every beer by contribution per tank-day, not by volume.**
   Your bestseller can be your worst earner. You will watch the ranking invert.

2. **Test the plan against physical capacity before you commit to it.**
   Taps, cold slots, and who is actually behind the bar on a Tuesday.

3. **Hand the model the constraints it cannot see.**
   Your margin structure, the Indian advertising rules, and your own floor.

Three outcomes, three use cases, one brewery and one decision running through all of them.

THE TRAIL YOU ARE ABOUT TO FOLLOW

Here is where we end up.

01

The ranking

Seven beers, ranked four ways. The order flips on the fourth.

02

The capacity check

Taps, cold slots and roster, run against the plan.

03

The compliance screen

Margin and Indian ad rules over the content plan.

Each one is a calculator you can rerun tonight on the take-home data, then on your own export tomorrow.

Nothing here is a demo you cannot repeat.

IN YOUR HAND RIGHT NOW

The prompt structure. Four lines.

1. The role and the decision	"You are helping me plan next month's production at a craft brewery."
2. The evidence, attached	"Attached are 18 months of taproom sales, brew records and the beer list."
3. The shape you want back	"Give me a table: beer, batches, litres, reason. Ranked. No commentary."
4. The honesty clause	"Tell me which numbers are thin instead of guessing."

Paste it, swap in your own file, and it works. Every prompt today is these four lines in a different order.

AI is not an
answer engine.

It is an evidence engine.

Everything in the next fifty minutes turns on that.

THE JOURNEY

From chasing numbers to reading evidence.

HOW IT WAS DONE	WHAT CHANGES
Three days pulling POS exports into one sheet	One thread, files attached, an answer in twenty minutes
The ranking you have always run, by volume	Four rankings side by side, including the one nobody runs
A plan that looks right until the Tuesday it breaks	The plan checked against tanks, taps and roster before you commit
An instinct you cannot show the bank	The same instinct, with the arithmetic attached

The instinct was never the problem. The evidence behind it took three days to assemble, so it usually never got assembled.

It is not looking anything up.

It is predicting the next word from everything it has read.

A malt that sounds right is, to the model, right.

It has read every brewing book ever published. It has never tasted a beer, and it has never seen your tank chart.

Fixable in one move. Stop asking it to remember. Start making it read.

Attach the evidence.
Ask for the shape.

That is the whole method. The four lines in your hand do exactly this.

Four pillars. Three of them need money.

Pillar 1 Time series

Tank logs read as a curve. Gravity, temperature and pressure over time, so a stuck ferment is flagged on day two instead of day five.

NEEDS LOGGERS ON THE TANKS

Pillar 2 Computer vision

The phone camera you already carry. Fill levels, label alignment, keg and line cleanliness, checked from photographs rather than a machine vision line.

NEEDS SOMEBODY TO SHOOT IT CONSISTENTLY

Pillar 3 Knowledge-based AI

The rules that never change, written down once. Yeast generation limits, tank turnaround, CIP intervals and style specs, checked against every plan.

NEEDS YOUR RULES OUT OF YOUR HEAD

Pillar 4 Generative AI

The brewer's copilot. Reads the exports you already have, explains the deviation, drafts the plan, and shows the arithmetic under it.

NEEDS A FILE AND FORTY MINUTES

Today is pillar four, on files you already export. The other three are capital projects, and none of them pay until this one does.

THE BOUNDARIES

What this session is, and is not.

WHAT IT IS	WHAT IT IS NOT
Reading your own numbers with a model	Training or building any model
Decisions you already make every month	Writing code, or a project for your IT person
Files you can export from your POS today	A new system to buy, install or migrate to
A method you can repeat on Wednesday	A vendor pitch. Nothing is sold from this stage

If you can export a sales report, you are qualified for everything in the next forty minutes.

USE CASE 1

What do we brew next?

The decision every brewery makes monthly, on instinct and a sales report.

10 minutes

WHAT THE MODEL SAYS

It is not wrong.

Deccan Lager is the volume leader.

73,249 litres over 18 months.

Brew more lager.

It is genuinely the bestseller, it has the largest unmet demand in the stockout file at 11,922 litres, and every sales signal points the same way.

THE OVERRULE

The sales file is right. The plan is wrong.

1 **66 litres per tank-day. Wheat does 158.**

1,000 litres of lager holds a fermenter for 15.1 days. Wheat, 6.3.

2 **Only three of eight tanks can even do it.**

FV1, FV2 and FV3 are the only fermenters with the glycol for lagering.

3 **The yeast is tired.**

Ten lager batches at generation 7 or 8, each held about three extra days.

The reason this plan is wrong is not in the sales file at all.

RANK IT A FOURTH WAY

And the answer inverts.

BEER	LITRES	GM PER LITRE	CONTRIBUTION PER TANK-DAY
Deccan Lager	73,249	INR 476	INR 34,000
Sahyadri Wheat	68,110	INR 522	INR 87,000
Ghat Road IPA	49,798	INR 588	INR 73,500

#1 of 7 by volume. #7 of 7 on contribution per tank-day. Wheat earns 2.6x per day of tank time.

Your bestseller is your worst earner per day of tank time.
And tank days are the only thing here you cannot buy more of this month.

TAKEAWAY 1 OF 3

Rank by contribution per tank-day, not by volume.

Tank days are the scarce resource. Rank by the thing you cannot buy more of.

USE CASE 2

Can you actually pour it?

The plan is now correct. Correct is not the same as possible.

8 minutes

WHAT THE MODEL SAYS

It is not wrong.

Good. Run the wheat.

Put the freed taps to work.

Promote it.

Nothing in the sales data or the production plan contradicts a word of that.

THE OVERRULE

Now tell me where the kegs go.

1 8 tap lines need 16 cold slots. The cold room holds 14.

Two lines always run with no cold backup keg, whatever the plan says.

2 232 days ran on one person behind the bar.

78 Tuesdays, 77 Wednesdays, 77 Mondays, across 18 months.

3 Every one of those 232 days already had more covers than it could serve.

Worst was 20 October 2025: 344 covers, against a roster that could serve 42.

The plan is right and it still fails on a Tuesday.

TAKEAWAY 2 OF 3

Test the plan against physical capacity before you commit.

Taps, cold slots and the roster kill more good plans than bad forecasting does.

USE CASE 3

How do you sell it?

The content that fills the room, against the content you are allowed to run.

8 minutes

WHAT THE MODEL SAYS

It is not wrong.

Price offer posts win.
11.4% engagement against 4.4%
for behind-the-scenes.

That is a real pattern, correctly found, in the data the model was given.

Two reasons. The second is the one nobody expects.

1 **They destroy the thing you are protecting.**

Contribution per litre: INR 536 at full price, INR 418 at 20% off, INR 363 at 30% off. The posts that fill the room fill it with your least valuable litres.

2 **You are not allowed to run them.**

67 posts in this history promote alcohol pricing directly. Indian surrogate advertising rules do not permit that. The best-performing content in the file should never have been published.

The model had no way to know the winning tactic is illegal here.

That is not a model problem. That is why you are in the room.

TAKEAWAY 3 OF 3

Hand it the constraints it cannot see.

Your margin structure, the advertising rules, and what you know about your own floor.

REPLAY THE TRAIL

Three correct analyses. Three wrong decisions.

01	What to brew	Sales said more lager. Tank time said it is the worst use of a fermenter
02	Can you pour it	The corrected plan still broke on a Tuesday, on cold slots and one bartender
03	How to sell it	The best-engaging posts were margin-destroying and not compliant

Not one of them was a data quality problem.

The machine assembles
the evidence.
You still decide.

FIND YOUR OWN

Six questions. The answers are the prompt.

CONTEXTUAL

Tell me how you do this task today.

PAIN

What is the hardest part of it, and why is it hard?

FREQUENCY

How often do you have to do it?

VALUE

Why does it matter to the company that it gets done?

EXISTING SOLUTION

What do you already do to solve it for yourself?

DRILL DOWN

What do you mean by that? Tell me more. Why does that matter?

Answer all six about one job in your own week, and you have written the prompt. That is the whole method, pointed at you.

QUIZ

One question. The malt order.

Your supplier runs the purchase history and calls you. Pilsner malt: 21,250 kg bought over 18 months, 170 kg left, 45-day lead time.

It is the most confident forecast in the whole dataset. So what is wrong with it?

THE ANSWER

61% of that Pilsner malt went into Deccan Lager, and this room cut the lager thirty minutes ago. The most predictable line in the forecast is about to collapse. Wheat malt runs the same 45-day lead time, so the replacement cannot be ordered reactively either.

WHAT YOU LEAVE WITH

All of it, free, today.

The prompt pack	Every prompt from the three use cases, in plain language
The four-line structure	The card in your hand, and why each line is there
The dataset	This brewery, 15 files. Rerun every demo before you risk real data
The data security page	The do and do not list, one page, on your fridge
Two bonus use cases	The malt forecast from the quiz, and the credit-burn habits
A free clinic	Bring your own export. No charge, no sales call

Nothing was sold from this stage, and nothing will be.

Questions.

Ten minutes, and the honest ones are the useful ones.

Ankur Napa
Growth Officer, Disruptive Advantage. Master Brewer turned data architect.
napaankur@gmail.com   linkedin.com/in/ankur-napa